

Assignment Project Exam Help

XJCO3221 Parallel Computation

<https://powcoder.com>

Peter Jimack

University of Leeds

Add WeChat powcoder

Lecture 5: Loop parallelism and data races

In Lecture 3 we saw how two problems could be parallelised:

- ➊ **Vector addition**, where two one-dimensional arrays were added together element-by-element.
- ➋ **Mandelbrot set**, where pixel colours were calculated to generate a 2D image.

<https://powcoder.com>

However, neither of these problems have any **data dependencies**.

- ➊ Each vector element was calculated **independently** of other elements.
- ➋ Each pixel colour was calculated **independently** of the others.

# Assignment Project Exam Help

Today we will look how to parallelise problems that **do** have data dependencies.

- Can lead to **data races** on shared memory systems.
- Behaviour then becomes **non-deterministic**.
- May require algorithm re-design to remove dependencies.

<https://powcoder.com>

## Add WeChat powcoder

We will then look at three examples of **loop parallel** problems and how their dependencies can be resolved.

## Example of a data race

Consider the following pseudocode<sup>1</sup> for two concurrent threads, where each thread accesses the same variable  $x$ .  $x=0$  at the start of the code segment.

Thread 0:

```
a = 1.0;  
a += 1.0;  
x = a;
```

Thread 1:

```
b = x;  
b += 2.0;  
x = b;
```

What value does  $x$  take at the end?

---

<sup>1</sup>From §2.6 of McCool *et al.*, *Structured parallel programming* (Morgan-Kaufman, 2012).

The result may differ each time the program is executed.

- An example of non-determinism [Lecture 5]

The scheduler runs threads depending on various factors, including other applications and OS tasks.

- Cannot predict which thread is **launched** first.
- The OS may **suspend** threads to make way for other tasks (*pre-emptive multitasking*).
- The instructions may become **interleaved**.

## Interleaved instructions (example)

Recall  $x=0$  initially.

```
1  a = x;    // Thread 0: a now 0
2  a = 1;    // Thread 0: a now 1
3  b = x;    // Thread 1: b now 0
4  x = a;    // Thread 0: x now 1
5  b = 2;    // Thread 1: b now 2
6  x = b;    // Thread 1: x now 2
```

# Add WeChat powcoder

In this example,  $x=2$  at the end.

- Possible to get  $x=1$  or  $x=3$  by different interleaving of instructions (*check left as exercise*).

This is known as a **data race** or a **race condition**.

- Result of calculation depends on which thread reaches its instructions first.
- **Non-deterministic**, which is usually undesirable.

<https://powcoder.com>

Only an issue for **shared memory**.

- If each thread had its own  $x$ , there would be no race.
- In this example, if each thread had its own  $x$ ,  $x=1$  for thread 0 and  $x=2$  for thread 1 at the end, regardless of any interleaving.

## Read-only does not lead to a data race

For a race condition to arise, **at least one** thread must **write** to `x`.

- No race if all threads just **read** `x`.

# Assignment Project Exam Help

There is no race for the following example, as both threads only read `x`:

<https://powcoder.com>

Thread 0:

Thread 1:

```
a = 2;  
a += 1.0;
```

```
b = x;  
b += 2.0;
```

# Add WeChat powcoder

For this example, `x=0` (and `a=1` and `b=2`) at the end.



Have assumed each thread executes its instructions in order.

- i.e. have assumed **sequential consistency**.

Modern compilers often rearrange instructions to improve performance.

- e.g. bring forward memory accesses, combine operations *etc.*
- The result is the same **in serial**.
- However, **multithreading** can confuse compilers.
- Can lead to unexpected results!

# The volatile keyword

You may read that the way to solve this is to declare variables as `volatile` (in C/C++). However, this is only partially correct<sup>1</sup>.

- `volatile` is for **special memory** such as that read/written by external device (*memory-mapped I/O*).
- The way compilers handle such variables is **not guaranteed** to work for multithreading.

<https://powcoder.com>

If this might be an issue, should use features **specific** to concurrent programming.

- e.g. memory fences, `std::atomic<>` in C++11 etc.
- Will come to atomics next lecture and Lecture 18.

---

<sup>1</sup>S. Meyers, *Effective modern C++* (O'Reilly, 2015).

Often we are required to parallelise **loops**.

- Known as **loop parallelism**.

- If there are **data dependencies**, may have implicit **data races** within the loop.

<https://powcoder.com>

There is no **systematic** way to parallelise loops.

- How to remove a dependency depends on context.
- Consider extra resources or loops to reach the correct solution.

Add WeChat powcoder

For the remainder of this lecture, will give examples of loops with data dependencies, and how to overcome them.

## Example 1: Redundant variable

Consider the following serial code:

```
1 float temp, a[n], b[n], c[n];  
2 ... // Initialise arrays b and c  
3  
4 int i;  
5 for( i=0; i<n; i++)  
6 {  
7     temp = 0.5f*( b[i] + c[i] );  
8     a[i] = temp;  
9 }
```

<https://powcoder.com>

Add WeChat powcoder

Here, temp is being used as a temporary variable.

- Sometimes useful to make (more complex) code easier to read.

Need to make temp a private (or local) variable:

```
1 #pragma omp parallel for
2 for( i=0; i<n; i++ )
3 {
4     float temp = 0.5f*( b[i] + c[i] );
5     a[i] = temp;
6 }
```

Can also use OpenMP's private clause

```
1 #pragma omp parallel for private(temp)
2 for( i=0; i<n; i++ )
3 {
4     temp = 0.5f*( b[i] + c[i] );
5     a[i] = temp;
6 }
```

cf. the inner loop counter in Lecture 3's Mandelbrot set example.

## Example 2: Shift dependency

Code on Minerva: `shiftDependency.c`

Consider a shift dependency:

```
1 float a[n];  
2 ... // initialise array a  
3  
4 int i;  
5 for( i=0; i<n-1; i++ )  
6   a[i] = a[i+1];
```

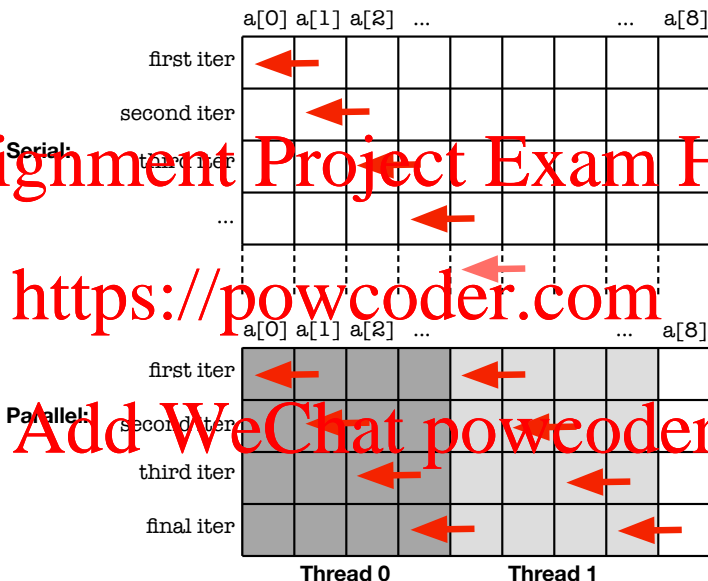
Naive parallelisation does not *quite* work:

```
1 #pragma omp parallel for  
2 for( i=0; i<n-1; i++ )  
3   a[i] = a[i+1];
```

# Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



A solution here is to **copy** the array **a** **before** the loop:

```
1 float atemp[n];  
2  
3 #pragma omp parallel for  
4 for( i=1, i<n, i++)  
5     atemp[i] = a[i];  
6  
7 #pragma omp parallel for  
8 for( i=0, i<n-1, i++)  
9     a[i] = atemp[i+1];
```

This comes at the expense of **additional resources**:

- **Memory** for the array **atemp**.
- **CPU time** to copy **a** to **atemp**.

Examples of **parallel overheads**.



## Example 3: Red-black Gauss-Seidel

Code on Minerva: redBlackGaussSeidel.c

Finally, consider this:

```
1 for( i=1; i<n-1; i++ )  
2   a[i] = 0.5f * ( a[i+1] + a[i-1] );
```

<https://powcoder.com>

Each element takes the average of the elements either side of it.

- Can be used to **smooth** vector **a**, e.g. in image transformation (blurring) with a 2D nested loop over pixels.
- Also arises in numerical computation - the **diffusion** or **heat equation** solved using the **Gauss-Seidel** method.

We **could** make a copy atemp as before.

- In numerical computation, this is the **Jacobi method**.

# Assignment Project Exam Help

However, this is undesirable in some situations:

- Typically repeat the loop until 'reaching' the solution, which is faster (*takes fewer iterations*) when not using a copy

<https://powcoder.com>

Instead we consider a **modified serial** variant which is more amenable to parallelisation.

Add WeChat powcoder

- Known as **red-black Gauss-Seidel**, as it bears a resemblance to red and black squares on a chessboard (in 2D).

Update **even** elements first, then **odd** elements<sup>1</sup>:

```
1 int redBlack;
2 for( redBlack=0; redBlack<2; redBlack++ )
3     for( i=1; i<n-1; i++ )
4         if( i%2 == redBlack )
5             {
6                 a[i] = 0.5f * ( a[i-1] + a[i+1] );
7             }
```

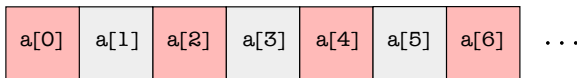
<https://powcoder.com>

The outer redBlack loop only loops over 2 values, 0 or 1.

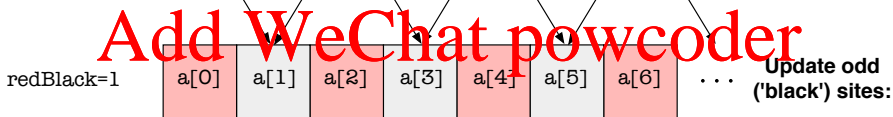
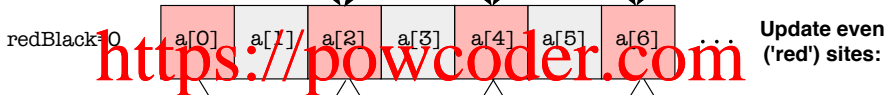
- When redBlack=0, only the elements of a[i] with i **even** are updated.
- Similarly, only the **odd** are updated when redBlack=1.

---

<sup>1</sup>Recall  $i\%2$  gives the remainder of division by 2, so e.g.  $i\%2==0$  if  $i$  is even.



Assignment Project Exam Help



Add WeChat powcoder

Note that for each loop, the calculations are **now independent**.

- We have **removed the dependency**, albeit by slightly changing the serial algorithm.

# Assignment Project Exam Help

Now clear how to parallelise:

```
1 for( redBlack=0; redBlack<2; redBlack++ )
2   #pragma omp parallel for
3   for( i=1; i<n-1; i++ )
4     if( i%2 == redBlack )
5     {
6       a[i] = 0.5f * ( a[i-1] + a[i+1] );
7     }
```

<https://powcoder.com>

Add WeChat powcoder

There are no dependencies **within** the i-loop, because  $a[i-1]$  and  $a[i+1]$  were/will be updated in the other redBlack loop.

# The 'best' parallel algorithm?

Notice that we changed Gauss-Seidel to the red-black variant to make it more efficient in parallel.

- For Gauss-Seidel, this is the standard parallel variant.

Perfectly acceptable to do this if <https://powcoder.com>

- 1 The solution reached is still 'correct.'
- 2 Parallel scaling is good (for large arrays).

## Add WeChat powcoder

As a general observation, the 'best' parallel algorithm need **not** be directly related to the 'best' serial algorithm.

Today we have seen how multiple threads with at least one writing to shared memory can lead to **data races**.

- Outcome is not predictable (*non-deterministic*).
- Can be a challenge to remove **data dependencies** from loops.
- May require additional resources not present in the serial version, i.e. **parallel overheads**.

# Add WeChat powcoder

Next time we will look at a way to **synchronise** threads within a parallel program and apply it to a linked list.