

Assignment Project Exam Help

XJCO3221 Parallel Computation

<https://powcoder.com>

Peter Jimack

University of Leeds
Add WeChat powcoder

Lecture 6: Critical regions and atomics

Previous lecture

Assignment Project Exam Help

In the last lecture we looked at **data races** and **loop parallelism**:

- If different threads access the same memory location, with at least one having write access, there is a potential **race condition**.
- Leads to **non-deterministic** behaviour.
- Can arise in loops as **data dependencies**.
- Often possible to remove these dependencies, sometimes at the expense of increased **parallel overheads**.

This lecture

Assignment Project Exam Help

In this lecture we will look at an important concept in parallel programming: **synchronisation**.

- Represents some form of **coordination between processing units** (threads *etc.*).
- Will arise in many contexts throughout this module.
- Briefly mentioned in Lecture 4 in the context of **fork-join**.

Now we will focus on using synchronisation to avoid data races.

- Define **critical regions** which can only be accessed by one thread at a time.
- **Atomics**: Critical regions specialised to single operations.

We will say more about atomics in Lecture 18.


Singly linked lists

Serial code on Minerva: `linkedList.c`

Assignment Project Exam Help

Linked lists are a form of **container** in which each item is linked together in a **chain**:

- First item is the **head**, with a global pointer `item_t *head`.
- Last item is the **tail**, with `next = NULL`.



Note this is a *singly* linked list - a **doubly** linked list has arrows 'both ways,' i.e. a field `item_t *prev;`.

List storage

Assignment Project Exam Help

For today's example, the data for each item is just a single integer.

To link into a list, convenient to use a struct in C:

```
1 typedef struct item {  
2     int data;  
3     struct item *next; // Next in the list  
4 } item_t;  
5  
6 item_t *head = NULL; // First item
```

New items are added using addToList:

```
1 for( i=0; i<numAdd; i++ )  
2     addToList( i );
```

Implementation of addToList()

```
1 void addToList( int data )
2 {
3     item_t *newItem = (item_t) malloc(sizeof(item_t));
4     newItem->data = data;
5     newItem->next = NULL;
6
7     // Find tail and add new item to it.
8     if( head==NULL ) // i.e. list is empty.
9     {
10         head = newItem;
11     }
12     else // Find end of list and add to it.
13     {
14         item_t *tail = head;
15         while( tail->next != NULL ) tail = tail->next;
16         tail->next = newItem;
17     }
18 }
```

Linked lists in parallel

Assignment Project Exam Help

Suppose we want to add multiple items in parallel for speed.

The obvious thing is to create multiple threads, and have each thread add items simultaneously.

In OpenMP, this could be achieved as follows:

```
1 #pragma omp parallel for
2 for( i=0; i<numAdd; i++ )
3     addToList( i )
```

- Multiple threads created at start of loop (*'fork'*).
- Each thread calls `addToList` some fraction of `numAdd` times.
- Check with `printList()` after the loop is complete (*'join'*).

Failure of `addToList()` when called in parallel

The items will not be added in the same order as in serial, but this may not be a problem as long as they are *somewhere* on the list.

Some items are being **lost**, especially when many are added.

- May differ run by run, i.e. non-deterministic.

Also, the memory allocated for lost items is never reclaimed.

- This **memory leak** could cause problems if the application was run for a long time (e.g. server).

The implementation of `addToList()` does not work in a **multithreaded context**. We say it is not **thread safe**.

Thread safety

A routine, library class, etc. is called **thread safe** if it works 'as advertised' in a **multithreaded** context.

If an API specification does not state whether or not a routine is thread safe, assume that it is not.

Note that being thread safe does **not** necessarily mean it is **efficient** in parallel.

- May have used a 'lazy' method to make a routine thread safe, but very slow.
- In this case may need to find an alternative that *does* scale in parallel, or develop your own solution.

The need for synchronisation

Assignment Project Exam Help

if `head` is global and therefore shared between threads.

This generates **data races** when adding a new item:

- 1 Multiple threads can have '`if(head==NULL)`' evaluate as true. Only one (the last) will become the new head.
- 2 When traversing the list, multiple threads may reach the tail at the same time. Again, only one is added.
- 3 A thread can read '`tail->next!=NULL`' fractionally before another thread sets '`tail->next`' to its item `t`.

If also *removing* (or 'popping') items, similar considerations would apply, although things would be more complicated.

Critical regions

Assignment Project Exam Help

It is not easy to remove the data dependencies in this case.

- Cannot make head local/private (*unique to the entire list*).
- List traversal is a while-loop (*trip count not known at start*).

<https://powcoder.com>

An alternative strategy is to ensure only one thread can access **critical regions** of code at a time.

- Implemented in OpenMP as `#pragma omp critical`
- Called **lock synchronisation**, for reasons that will become clear next lecture.

Add WeChat powcoder

```
#pragma omp critical
```

Assignment Project Exam Help

Critical region

Only one thread is allowed to be in a critical region at a time.
Until it leaves, **no other threads** are allowed to enter.

<https://powcoder.com>

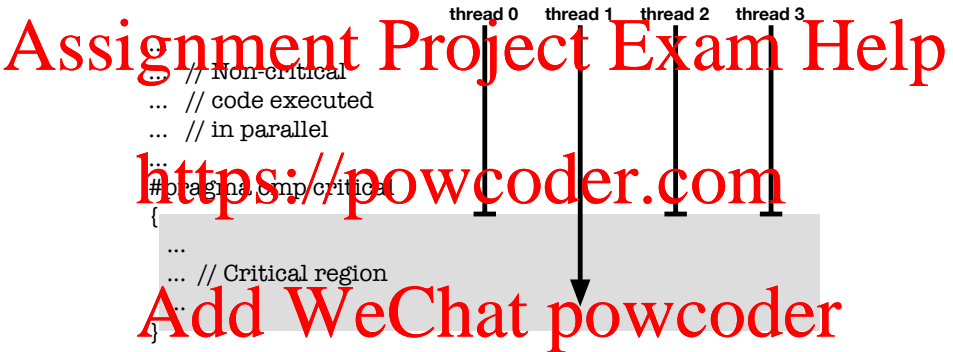
A **critical region** is defined in OpenMP as follows:

```
1 #pragma omp critical
2 {
3     ... // Critical region
4 }
```

Add WeChat powcoder

The critical region is defined by the **scope**, *i.e.* from '{' to '}'.

Example for 4 threads



- Thread 1 reaches the critical region **first**.
- **No other thread** can enter until it leaves.
- **Exactly one** thread may then enter.

Performance

Assignment Project Exam Help

There can be a significant performance penalty for critical regions:

- 1 Need some mechanism to **synchronise** the threads on entering and leaving the region [*cf. Lecture 7*]
- 2 Threads 'blocked' at the start will be idle. This leads to poor **load balancing** [*cf. Lecture 13*].
- 3 The scheduler may **suspend** idle threads in favour of another (not necessarily yours!). Suspension and restart incur penalties.

Serialisation

Since only one thread can be in a critical region at any time, the critical code is **executed in serial**.

This is known as **serialisation**.

<https://powcoder.com>
Amdahl's law and the **Gustafson-Barsis law** from Lecture 4:

- Maximum speed-up S in terms of the **serial fraction** f .
- By serialising regions of code we are **increasing** the value of f .
- The maximum speed-up S is **reduced**, especially for Amdahl's law (*i.e.* strong scaling), which predicts $S \leq 1/f$.

It is therefore important to restrict the **number** and **size** of critical regions to ensure reasonable parallel performance.

First attempt: Serialise calls to addToList()

```
1 #pragma omp parallel for  
2 for( i=0; i<numAdd; i++ )  
3     #pragma omp critical  
4     {  
5         addToList( i );  
6     }
```

Assignment Project Exam Help

<https://powcoder.com>

This works, but parallel performance is poor.

- Essentially the whole loop has been serialised
- Would be better off leaving it in serial, *i.e.*

```
1 for( i=0; i<numAdd; i++ )  
2     addToList( i );
```

Add WeChat powcoder

Attempt 2: Serialise list traversal

Note that the start of `addToList()` has no data dependences

```
1 item_t *newItem = (item_t*) malloc(sizeof(item_t));  
2 newItem->data = data;  
3 newItem->next = NULL;
```

- `newItem` is local/private to each thread.
- Each thread will create its own item **independently of other threads**.
- Each value of loop counter `i` will be mapped one to one to a value of `data`.

This is the behaviour we want! (so far...)

The data dependencies in the remainder of addToList() can be removed by placing this portion in a critical region:

```
1 #pragma omp critical
2 {
3     if( head==NULL )
4     {
5         head = newItem;
6     }
7     else
8     {
9         item = *tail = head;
10        while (tail->next != NULL) tail = tail->next;
11        tail->next = newItem;
12    }
13 }
```

Performance *slightly* improved compared to the previous attempt.

Making routines thread safe

Assignment Project Exam Help

Note the strategy followed for making addToList() **thread safe**:

- 1 **Identify** data dependencies.
- 2 Enclose in **critical regions**.
- 3 **Reduce** size and/or number of critical regions until required performance achieved.
- 4 If further scaling benefits required, may need to rethink the algorithm completely.

<https://powcoder.com>

Add WeChat powcoder

Atomics

Assignment Project Exam Help

Often the data dependency is only a single arithmetic operation.

For instance, counting the number of items in an array of size n that obey some condition:

```
1 int count = 0;
2 for( i=0; i<n; i++)
3 {
4     ...
5     if( condition[i] ) count++;
6 }
```

<https://powcoder.com>

Add WeChat powcoder

The command `count++` is a **data race**:

- Two threads may read the **old** value of `count` simultaneously.
- New `count` may not be the old value $+2$ [*cf. Lecture 5*].

Critical region

Assignment Project Exam Help

Using a critical region will work ...

```
1 int count = 0;
2 #pragma omp parallel for
3 for( i=0; i<n; i++ )
4 {
5     ...
6     if( condition[i] )
7     #pragma omp critical
8     {
9         count++;
10    }
11 }
```

<https://powcoder.com>

Add WeChat powcoder

... but has the usual overheads of a critical region.

Atomic instructions

Assignment Project Exam Help

Because this is a common situation, most compilers/hardware can perform the necessary synchronisation efficiently.

- **Very low overhead.**

<https://powcoder.com>

These are known as **atomic instructions**.

Add WeChat powcoder

Can think of atomics as a special type of critical region specialised to single arithmetic operations that exploits hardware primitives.

GPUs also support atomic instructions; we will look at these more closely in Lecture 18.

Atomics in OpenMP

Assignment Project Exam Help

Atomics are implemented in OpenMP as follows:

```
1 int count = 0;
2 #pragma omp parallel for
3 for( i=0; i<n; i++ )
4 {
5     ...
6     if( condition[i] )
7         #pragma omp atomic
8         count++;
9 }
```

<https://powcoder.com>

Add WeChat powcoder

Note there is no scope ('{'...'}') after `#pragma omp atomic`.

- Only works on single instructions.

Summary and next lecture

Assignment Project Exam Help

Today we have looked at using **critical regions**:

- Can avoid data races by **serialising** blocks of code.
- Corresponding performance loss.
- Reduce overhead for single arithmetic instructions by using **atomics**.

<https://powcoder.com>

Add WeChat powcoder

Next lecture we will look in more detail at how this synchronisation is achieved.