

Assignment Project Exam Help

XJCO3221 Parallel Computation

<https://powcoder.com>

Peter Jimack

University of Leeds
Add WeChat powcoder

Lecture 7: Lock and mutexes

Previous lecture

Assignment Project Exam Help

In the last lecture we saw how critical regions of code could be **serialised**:

- Only one thread can enter the region at a time, a form of **coordination**.
- Avoids **data races**.
- Can incur a significant **performance penalty**.
- Implemented in OpenMP as `#pragma omp critical`.
- Single arithmetic instructions can be optimised by using **atomic** instructions (`#pragma omp atomic`).

Today's lecture

For today's final lecture on shared memory parallelism, we will look at what is going on 'behind the scenes'.

- Thread coordination performed using **locks**, sometimes known as **mutexes**.
- Locks can control access to **data structures**.
- Multiple locks can improve **performance** of memory access.
- However, multiple locks can give rise to **deadlock**.

<https://powcoder.com>
Add WeChat powcoder

This lecture is largely theoretical¹ and will **not** help with any courseworks, but the material may appear in the exam.

¹There **are** code examples for this lecture, and a question on Worksheet 1.

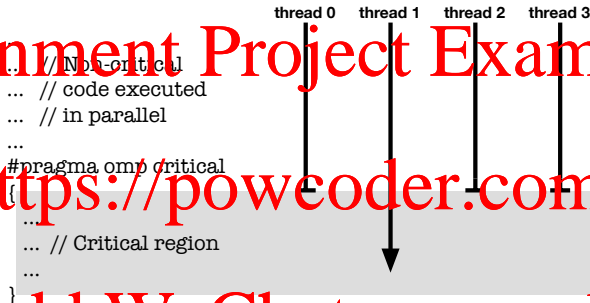
Recap: Critical regions

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Instructions before `#pragma omp critical` executed **concurrently** (e.g. if in a parallel loop).
- Instructions in the scope (`{` to `}`) only executed by **one thread at a time**.
- Other threads blocked from entering; they are **idle**.



Thread coordination with locks

Assignment Project Exam Help

This **synchronisation** is performed using a **lock**:

- Single lock for the critical region.
- Can be in one of two states: **locked** and **unlocked**.
- The first thread to reach the critical region **locks** it.
 - Also known as **acquiring the lock**.
 - This thread is said to be the lock's **owner**.
- No other threads can enter the region ('acquire the lock') until it becomes unlocked.
- The owning thread **unlocks** (or **releases**) it when leaving the region, allowing another thread to take over ownership.

Critical region using a lock

OpenMP:

```
1 // Multiple threads executing concurrently,  
2 // (e.g. parallel loop).  
3  
4  
5 #pragma omp critical  
6 {  
7     ...  
8     ... // Critical code.  
9     ...  
10 }
```

Lock pseudocode:

```
1 // All threads access a  
2 // single lock object.  
3 lock_t regionLock;  
4  
5 regionLock.lock();  
6  
7 ...  
8 ... // Critical code.  
9 ...  
10 regionLock.unlock();
```

`regionLock.lock()` does not return until the thread has **acquired** the lock; it is said to be **blocking**.

Implementations of locks

Assignment Project Exam Help

Most parallel APIs support **locks**, although they are sometimes called **mutexes** as they control m**utual** e**xclusion**:

- Java's Lock interface (in `java.util.concurrent.locks`).
- `std::mutex` in C++11 (in `<mutex>`).
- `pthread_mutex_t` in the pthreads library (C/C++).

Add WeChat powcoder

When implemented as classes, they are typically **opaque**:

- The user does not have access to instance variables or details of the implementation.

Locks in OpenMP

OpenMP runtime library also supports locks:

```
1 #include <omp.h>
2
3 // Initialise lock (opaquely).
4 omp_lock_t regionLock;
5 omp_init_lock(&regionLock);
6 ... // (in parallel).
7 omp_set_lock(&regionLock); // LOCK.
8 ... // (critical code).
9 omp_unset_lock(&regionLock); // UNLOCK.
10 ...
11 // Deallocate the lock.
12 omp_destroy_lock(&regionLock);
```

You *could* implement your own critical region this way, although it is easier to use `#pragma omp critical`.

Programming locks

Assignment Project Exam Help

Note there is no explicit link between the lock region, lock and the critical region of code, or data structure, that it is trying to protect.

<https://powcoder.com>

It is down to the **programmer** to correctly link each lock with its associated block of critical code, or data structure.

Add WeChat powcoder

This gives greater **flexibility**, but also greater scope for programming errors.

- Could use a struct or class to keep the lock with the data it is protecting, with the lock private/protected.

Lock mistakes (1): Forgetting to lock()

```
1 lock_t regionLock;  
2 ...  
3 //regionLock.lock(); // Forgot to lock()!  
4 ...  
5 ... // Critical code  
6 ...  
7 regionLock.unlock();
```

This is precisely the situation we were trying to avoid!

- All threads enter the critical region.
- **Race conditions** become a possibility.

unlock() will have no effect, except possibly a small performance overhead¹.

¹Generally, this depends on the API: In C++11, attempting to unlock a `std::mutex` that is **not** locked leads to undefined behaviour.

Lock mistakes (2): Forgetting to unlock()

Assignment Project Exam Help

```
1 lock_t regionLock;  
2 ...  
3 regionLock.lock();  
4 ...  
5 ... // Critical code  
6 ...  
7 //regionLock.unlock(); // Forgot to unlock!
```

<https://powcoder.com>

- The first thread **exclusively** enters the critical region.
- It never **releases** the lock.
- Therefore no other thread can **acquire** the lock.
- **All other threads remain idle at lock().**

RAII = Resource Acquisition Is Initialisation.

Assignment Project Exam Help

This second mistake is easier to make than it seems:

- The critical code may throw an **exception** (C++/Java).
- A break or continue command may jump over unlock().

<https://powcoder.com>

May be support for locks that automatically release when they leave their scope.

- If defined at start of a routine, automatically released at end of routine **however it reached there**.
- e.g. `std::lock_guard<std::mutex>` in C++11.

This mechanism is generally known as RAII, for Resource Acquisition Is Initialisation.

Multiple locks

Code on Minerva: `multipleLockCopy.c`

Assignment Project Exam Help

Suppose we want to copy randomly selected elements of an array of size N to another randomly-selected element.

- Decide to use a **lock** to control data writes.

```
1 #pragma omp parallel for
2 for( n=0; n<N; n++)
3 {
4     i = rand() % N;
5     j = rand() % N;
6
7     omp_set_lock( &entireLock );    // Lock.
8     data[j] = data[i];              // Safe copy.
9     omp_unset_lock( &entireLock );  // Unlock.
10 }
```

Multiple locks for memory access


This works, but is very inefficient.

- Only one thread can access the array data at a time.
- Even though only writing to one value.



Better to use **multiple locks** spanning the array:

- Different threads can write to different regions of the array **simultaneously**.
- Less **idle time** spent waiting for a lock to be released.



Using multiple locks is measurably faster (*try the code*):

```
1 omp_lock_t partialLocks[M];  
2  
3 // Initialise M locks near start of code.  
4 ...  
5 // Identify lock for writing to array.  
6 int lock = M*j/N;  
7 omp_set_lock( &partialLocks[lock] );  
8 data[j] = data[i];  
9 omp_unset_lock( &partialLocks[lock] );  
10 ...  
11 // Destroy all locks at end of code.
```

Note we only lock for the **write** to element j

- Recall that just reading does **not** invoke a data race.

Multiple locks for swapping

Code on Minerva: `multipleLockSwap.c`

Assignment Project Exam Help

Suppose now we want to swap elements i and j .

- Want to protect **each write** during the swap.

If access to the whole array was governed by a single lock, this would be straightforward to implement.

```
1 omp_set_lock( &entireLock );  
2  
3 // Writes to both data[i] and data[j].  
4 float temp = data[i];  
5 data[i] = data[j];  
6 data[j] = temp;  
7  
8 omp_unset_lock( &entireLock );
```

However, performance would again be poor.

Multiple locks for swapping

We might think of using two locks, one for each region of the array being written to:

```
1 int lock_i = M*i/N;  
2 int lock_j = M*j/N;  
3  
4 omp_set_lock( &partialLocks[lock_i] );  
5 omp_set_lock( &partialLocks[lock_j] );  
6  
7 float temp = data[i];  
8 data[i] = data[j];  
9 data[j] = temp;  
10  
11 omp_unset_lock( &partialLocks[lock_i] );  
12 omp_unset_lock( &partialLocks[lock_j] );
```

Try this out!

Why does this fail? Deadlock

Assignment Project Exam Help

Suppose one thread tries to lock `lock_i` then `lock_j` and **simultaneously** another tries to lock `lock_j` then `lock_i`.

- Thread 1 **owns** `lock_i`, **waits** for `lock_j` to be released.
- Thread 2 **owns** `lock_j`, **waits** for `lock_i` to be released.

Since each thread is waiting for the other lock, they will never release the lock they own. **They will both wait forever.**

Add WeChat powcoder

Threads waiting for synchronisation events that will never occur is known as **deadlock**.

The 'forgetting to `unlock()`' example earlier is also **deadlock**.

Nested critical regions

Code on Minerva: `nestedCriticalRegion.c`

Assignment Project Exam Help

Another problem is when `lock_i == lock_j`. A simpler example where this occurs is for **nested critical regions**:

```
1 // Outer critical region
2 omp_set_lock( &lock )
3
4 // Inner critical region
5 omp_set_lock( &lock );
6 ...
7 omp_unset_lock( &lock ); // End of inner region
8
9 omp_unset_lock( &lock ); // End of outer region
```

In OpenMP, this will also **deadlock**.

- A thread that **owns** a lock cannot **re-acquire** the lock.

Nested #pragma omp critical

Assignment Project Exam Help

OpenMP does not allow nested critical regions:

```
1 #pragma omp critical
2 {
3     ..
4     #pragma omp critical
5     {
6         ...
7     }
8 }
```

<https://powcoder.com>

Add WeChat powcoder

... will **not** compile.

- The **same** lock is being used by **both** critical sections.
- The same problem as in the previous slide.

Named critical regions

This can be resolved by using **named** critical regions:

```
1 #pragma omp critical (OUTER)
2 {
3     ...
4     #pragma omp critical (INNER)
5     {
6         ...
7     }
8 }
```

- OUTER and INNER are user-defined labels.
- Each unique label corresponds to a unique lock.
- You are implicitly using a different lock for each critical region, so no thread tries to re-acquire a lock it already owns.

Reacquiring locks

OpenMP code **deadlocks** if a thread tries to reacquire a lock it already owns.

- Just as it does not support nested critical regions.

This is primarily due to performance considerations.

- Would incur an execution overhead that would not always be necessary.

Not all parallel/concurrent APIs impose the same requirements - need to check the documentation!

- e.g. For C++11's `std::mutex`, the behaviour is undefined.
- Should also check documentation if attempting to `unlock` a lock that was **not** acquired.

Summary of shared memory systems

Assignment Project Exam Help

Lecture	Content	Key points
2	Architectures and OpenMP	Cache coherency; false sharing; kernel <i>versus</i> user threads.
3	Data parallel problems	Fork-join construct; non-determinism; embarrassingly parallel problems.
4	Theory	Amdahl's law (strong scaling); Gustafson-Barsis law (weak scaling).
5	Data races	Loop parallelism; data dependencies.
6	Critical regions	Thread coordination; thread safety; serialisation; atomics.
7	Locks/mutexes	Performance costs for locks; dead-lock; named critical regions.

Next lecture

Assignment Project Exam Help

Next time we will start to look at **distributed memory systems**:

- Multiple **processes** with their own **heap memory**.
- Examples: Clusters, supercomputers

<https://powcoder.com>

Not surprisingly, data races are *not* an issue, but many of the other aspects we have covered are:

- Non-determinism, scaling, deadlock, data and loop parallelism, . . .

Add WeChat powcoder