

Assignment Project Exam Help

XJCO3221 Parallel Computation

<https://powcoder.com>

Peter Jimack

University of Leeds
Add WeChat powcoder

Lecture 8: Introduction to distributed memory parallelism

Previous lectures

Assignment Project Exam Help

In the last six lectures we looked at **shared memory parallelism** (SMP) relevant to e.g. multi-core CPUs:

- Each **processing unit** (e.g. thread, core) sees **all** memory.
- Want to achieve good **scaling**, i.e. speed-up for increasing numbers of cores.
- Without proper **synchronisation**, results can be **non-deterministic**.
- Dependencies can lead to **data races**.
- Can reach **deadlock** if threads wait for synchronisation events that never occur.

<https://powcoder.com>

Add WeChat powcoder

This lecture

Assignment Project Exam Help

This lecture is the first of six on **distributed memory parallelism**, and we will see that some (but not all) of these issues remain relevant:


- Each **processing unit** sees only a **fraction** of total memory.
- **Data dependencies** treated using **explicit communication**.
 - No **data races**.
- Performance considerations remain the same except now the primary parallel overhead is **communication**.
- Improper synchronisation can still lead to **non-determinism** and **deadlock**.

Distributed memory systems

Assignment Project Exam Help

Multiple processes (rather than threads) that communicate via an interconnection network or 'interconnect'.

- For instance, one process per **node**, e.g. desktop machine.
- Each process has its own **heap memory**
- If a process needs data currently held on another node's memory, must **communicate** over the network.



Current fastest supercomputer¹

Fujitsu Fugaku, RIKEN, Kobe, Japan

- ARM-based A64FX CPU.
- 48 compute cores, and 2 or 4 assistant cores.
- Total 7,680,848 cores.
- No GPUs.
- Draws nearly 30MW of power.
- Benchmarked ≈ 442 PFLOPS.
- 1 PFLOPS = 10^{15} FLOPS.
- 1 FLOPS = 1 floating point operation per second.

¹As of Nov. 2021; top500.org.



Clusters as distributed systems

Assignment Project Exam Help
Supercomputers share features with other distributed systems such as data centres:

- Nodes perform calculations in parallel.
- Coordination requires explicit communication; there is no "global clock."
- May have high energy demand and cooling requirements.

Here focus on High Performance Computing (HPC) clusters:

- Individual cluster nodes use the same **operating system**.
- Cannot usually be **addressed individually**.
- Requires a special **job scheduler**.

The interconnection network or 'interconnect'

For the local area networks within HPC clusters, communication between nodes is carried over high performance **interconnects**.

- **Gigabit Ethernet** and **InfiniBand** are the most common¹.
- Latencies (*i.e.* delays) of around $1\mu\text{s}$.
- Bandwidths (*i.e.* throughput) of around 1-100 Gb/s.

These numbers are improving with time but **more slowly than CPU performance**

Add WeChat powcoder

The need to reduce communication overheads will only become more important in the foreseeable future.

¹As of Nov. 2021; see top500.org.

Network topology

If data sent via intermediary nodes, latency is increased

- Each node must parse data packet and decide where to send.


Therefore want smallest **paths** between nodes.

Network as a **graph** $G(V, E)$:

- V = nodes (*vertices*).
- E = connections (*edges*).

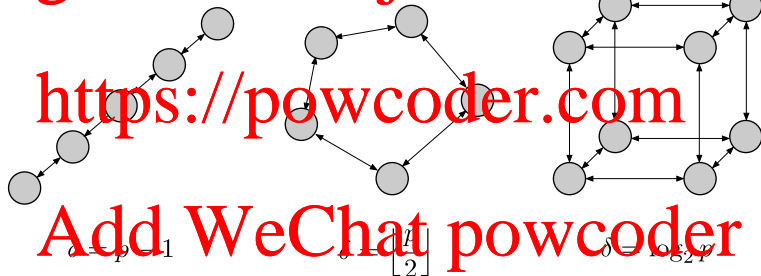
Want G with smallest **diameter** δ
(largest path length between nodes).

A **complete graph** (right) has $\delta = 1$,
but is impractical (*too many connections for each machine*).



Example topologies for p nodes

Assignment Project Exam Help



Hypercube topology preferred due to its short path lengths¹.

¹Rauber and Rünger, *Parallel programming for multicore and cluster systems* (Springer, 2013).

Processes *versus* threads

Recall from Lecture 2 that **processes** communicate with other processes using e.g. sockets.

- Must have **at least** one process per node to communicate across the network.

<https://powcoder.com>


For multi-core nodes, could have one **multi-threaded process** per node, with one thread per core.

- Avoids communication **within** a node.
- Combination of OpenMP and MPI is quite common ('*hybrid*').

For simplicity, we consider one **single-threaded process per core**, and therefore **multiple processes per node**.


Example for quad core nodes

One 4-thread process per node




node 1

single process




node 2

4 one-thread processes per node



node 1

4 processes



node 2

Books

Assignment Project Exam Help

Willkinson and Allen [Lecture 1] covers distributed memory parallelism (MPI), and a little OpenMP, but no GPU.

- General parallel algorithms but few code examples.
- Slightly old (2005) and covers architectures we will not consider (e.g. *distributed shared memory systems*).

A more practical book for MPI coding is:

- **Parallel Programming with MPI**, Pacheco (Morgan-Kaufman).
 - Old (1997), only covers distributed memory systems and MPI.
 - Many code examples and snippets.

Distributed HPC programming

Assignment Project Exam Help

For distributed HPC, there is essentially only one option¹ **MPI**

- Stands for **Message Passing Interface**.
- Specifies a **standard** for communication (*'message passing'*).
- MPI v1.0 finalised in 1994.
- MPI v3.0 finalised in 2012, now widely implemented.
- Fully supports C, C++ and FORTRAN.
- Most online examples are in one of these languages.
- Unofficial bindings for Java, MATLAB, Python, ...

¹Has superseded PVM = Parallel Virtual Machine (1989). Others such as Spark, Chapel etc. not (yet?) widely used in HPC.

Implementations

The MPI standard only defines the interface, it's still down to a vendor to provide an **implementation**.

- Code should be **portable** between implementations.

<https://powcoder.com>

There are various **freely available implementations**:

- **MPICH**: www.mpich.org
- **OpenMPI**: www.open-mpi.org
- Don't confuse OpenMPI with OpenMP!

Add WeChat powcoder

There are also commercial implementations:

- e.g. Intel MPI, Spectrum MPI (IBM).

Installing MPI

The system `cloud-hpc1.leeds.ac.uk` has OpenMPI¹ installed:

Assignment Project Exam Help

```
module load mpi/openmpi-x86_64
```

For personal Unix machines should be straightforward to install
(*cf. links on previous slide*).

- Mac users might like to try homebrew.

On Windows machines, Microsoft MPI² is free.

- Based on MPICH.

¹Note the linux command “`module avail`” shows what modules are installed.

²<https://docs.microsoft.com/en-us/message-passing-interface/microsoft-mpi>

Building an MPI program

Assignment Project Exam Help

Need to use a special compiler for MPI programs:

- Standard installation includes `mpicc`, `mpic++`, `mpifort`.
- Essentially a wrapper around a standard compiler.
- Passes command line arguments to the C compiler.

<https://powcoder.com>

For example, to compile a file `helloWorld.c`:

`mpicc -Wall -o helloWorld helloWorld.c`

- Will generate the executable `helloWorld`.
- All warnings on (`'-Wall'`).
- Add e.g. `-lm` for the maths library.

Add WeChat powcoder

Executing an MPI program

Also need a special **launcher** to execute an MPI program¹.

For multiple processes all on the same local machine:

```
mpiexec -n 2 ./helloWorld
```

- Creates 2 processes running the **same** program.
- Trying to launch more processes than cores may lead to an error (*'too many slots'*)².
- mpirun is the same/very similar to mpiexec.

Best to develop/debug code on a single machine (e.g. login node of cloud-hpc1.leeds.ac.uk), then run on multiple cores in batch mode for e.g. timing runs.

¹Executing as usual ('./helloWorld') will launch *one* process, i.e. serial.

²With OpenMPI, can override with the argument `-oversubscribe`.

Launching via the batch queue

The system `cloud.hpc1.leeds.ac.uk` has been set up to allow access to two 8-core nodes via `slurm`.

- Follow a similar approach to running batch jobs for OpenMP:

- `sbatch script.sh`

- Below is an example script...

```
#!/bin/bash
```

```
#Request a single node, and 8 cores (adjust as necessary)
```

```
#SBATCH -N1 -n8
```

```
module add mpi/openmpi3-x86_64
```

```
mpiexec -n 8 ./helloWorld
```

A 'Hello World' example

```
1 #include "stdio.h"
2 #include "stdlib.h"
3 #include "mpi.h" // Need to include mpi.h
4
5 int main( int argc, char **argv )
6 {
7     int numprocs, rank;
8
9     MPI_Init( &argc, &argv );
10    MPI_Comm_size( MPI_COMM_WORLD, &numprocs );
11    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
12
13    printf( "Process %d of %d.\n", rank, numprocs );
14
15    MPI_Finalize();
16    return EXIT_SUCCESS;
17 }
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Initialising and finalising

Assignment Project Exam Help

The first MPI call **must** be `MPI_Init()`:

- Pass command line arguments `argc` and `argv`.
- Will remove arguments relevant to MPI.
- Specific to the implementation and not of interest here.

<https://powcoder.com>

The final MPI call **must** be `MPI_Finalize()`:

- Note the JS spelling *finalize* not *finalise*.

Add WeChat powcoder

Any MPI calls before `MPI_Init()` or after `MPI_Finalize()` will result in a runtime error.

Number of processes and rank

Assignment Project Exam Help

`MPI_Comm_size(MPI_COMM_WORLD, &numprocs)`

- Sets numprocs to the **total** number of processes.
- Should return the '-n' argument in mpiexec.
- Similar to `omp_max_thread_num()`.

<https://powcoder.com>

`MPI_Comm_rank(MPI_COMM_WORLD, &rank)`

- Sets rank to the process number, known as the **rank** in MPI.
- Ranges from 0 to numprocs-1 inclusive.
- Similar to `omp_get_thread_num()`.

Add WeChat powcoder

Communicators

Assignment Project Exam Help

For our purposes, whenever you see an MPI call with the argument **communicator**, just use `MPI_COMM_WORLD`:

- Means 'all processes available to us.'
- The **only** communicator we consider in this course.

<https://powcoder.com>

In general, communicators allow processes to be **partitioned**.

- e.g. when developing a parallel library, don't want the library processes to accidentally communicate with application processes.
- An advanced feature we won't consider.

Add WeChat powcoder

Summary and next lecture

Assignment Project Exam Help

Today we have started looking at **distributed memory parallelism**:

- Realised in **clusters** and **supercomputers**.
- Requires **communication** between **nodes**.
- For HPC, use **MPI = Message Passing Interface**.
- Seen how to build and execute a 'Hello World' program.

Add WeChat powcoder

Next time we will see how MPI supports communication between processes, and use this to solve real problems.