

# Assignment Project Exam Help

XJCO3221 Parallel Computation

<https://powcoder.com>

Peter Jimack

University of Leeds  
Add WeChat powcoder

Lecture 11: Reduction

## Previous lectures

# Assignment Project Exam Help

In the last lecture we looked at **data reorganisation** and **collective communication**:

- **Communication** is usually the most significant overhead for distributed systems.
- **Collective communication** involves multiple processes in a one-to-many, many-to-one or many-to-many pattern.
- Reduce the communication time  $t_{comm}$  compared to a loop of point-to-point communications.
- In MPI: `MPI_Bcast()`, `MPI_Scatter()`, `MPI_Gather()`.

## Today's lecture

# Assignment Project Exam Help

Today we will look at a common combination of data reorganisation and computation: **Reduction**.

- Reduces a data set to one of a smaller size.
- Important for both shared and distributed memory systems.
- Support for many parallel APIs, including OpenMP and MPI.
- Often optimised using a **binary tree**.
- Binary trees also useful for collective communication.

<https://powcoder.com>

Add WeChat powcoder

## Reminder: Serial reduction

# Assignment Project Exam Help

- Start with a large data set.
- Apply **binary operations** to *reduce* to a smaller set.

Example 1: summing the elements of an array

```
1 sum = 0;  
2 for( i=0; i<N; i++ )  
3   sum += a[i];
```

Add WeChat powcoder

Example 2: Finding the maximum element

```
1 max = a[0];  
2 for( i=1; i<N; i++ )  
3   if( a[i]>max ) max = a[i];
```


Note each operation is performed **sequentially**.

Total computation time

$t_{\text{comp}}$  is proportional to the array size  $n$ .


- i.e. the **time complexity** is  $\mathcal{O}(n)$ .

If these were **processing units**, most would be **idle** throughout most of the calculation.



## Parallel reduction

Ideally we would want to perform all calculations simultaneously:



This *would* have a time complexity of  $t_{\text{comp}} = \mathcal{O}(1)$ , but is not possible to achieve in practice.

For now, note that:

**Any** parallel reduction **must** change the sequence of calculations

Some concrete examples will be given later in this lecture.

## Recap: Commutativity and associativity

# Assignment Project Exam Help

Let  $\otimes$  denote a general binary operator:  $x = a \otimes b$ .

As parallel reduction alters the sequence in which calculations are performed,  $\otimes$  must be **associative**:

<https://powcoder.com>  
An operator  $\otimes$  is **associative** if  $a \otimes (b \otimes c) = (a \otimes b) \otimes c$

If  $\otimes$  is only **approximately associative**, the result of parallel reduction will be (slightly) **different from serial reduction**.

# Add WeChat powcoder

Some parallel reduction algorithms also require  $\otimes$  to be **commutative**:

An operator  $\otimes$  is **commutative** if  $a \otimes b = b \otimes a$

## Commutativity and associativity (examples)

# Assignment Project Exam Help

Combination	Examples
Associative <b>and</b> commutative	max, min, Boolean AND, OR, XOR, <b>exact</b> addition and multiplication
Associative; <b>not</b> commutative	Matrix multiplication
Commutative; <b>not</b> associative	<b>Finite precision</b> floating point addition and multiplication <sup>1</sup> , signed saturated addition <sup>2</sup>
<b>Neither</b> commutative <b>nor</b> associative	Subtraction, division

<sup>1</sup>Only *approximately* associative. See Worksheet 2 Question 6.

<sup>2</sup>e.g.  $\text{fn}(a,b)=(a+b < 1 ? a+b : 1)$  with  $a=0.8$ ,  $b=0.5$  and  $c=-0.3$ .




# MapReduce

## Assignment Project Exam Help

An important application of reduction is as part of the **MapReduce** pattern<sup>1</sup>:

- Fusion of a **map** followed by a **reduction**.
- Can avoid the need for **synchronisation** after the map.



<sup>1</sup>McCool *et al.*, *Structured parallel programming* (Morgan-Kaufmann, 2012).

## Distributed systems example

# Assignment Project Exam Help

Suppose a database is distributed over nodes in a cluster.

- Each node has access to part of the full database.

Suppose a user initiates a search. We could use **MapReduce**:

- Each node searches its local database (*'map'*).
- Local results are combined to give the required global result (*'reduce'*).

This **MapReduce** was developed by Google and was one of the reasons for their early success.

## Example: Vector dot product

Consider the **vector dot** product (aka inner or scalar product):

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=1}^n a_i b_i$$

In serial<sup>1</sup>:

```
1 float dot=0.0;  
2 for( i=0; i<n; i++ )  
3   dot += a[i] * b[i];
```

Note this is a **map** (*the multiplication*) followed by a **reduction** (*the summation*).

---

<sup>1</sup>Recall maths indexing starts from 1 but computer indexing starts from 0.

## Reduction in OpenMP

Code on Minerva: dotProduct\_OpenMP.c

# Assignment Project Exam Help

In OpenMP (i.e. shared memory systems), reduction is supported by the **reduction clause**:

```
1 float dot=0.0;  
2 #pragma omp parallel for reduction(+:dot)  
3 for (i=0; i<N; i++)  
4     dot += a[i] * b[i];
```

- Specify the **binary operation** (‘+’) and the **target variable** (‘dot’).
- Compiler and runtime will implement an **efficient** reduction **for the given architecture**.
- Details of the implementation **opaque** to the user.

## Reduction in MPI

Code on Minerva: dotProduct\_MPI.c

# Assignment Project Exam Help

For MPI, distribute the full arrays of rank 0 to local arrays on each process using `MPI_Scatter()`<sup>1</sup>:

```
1 MPI_Scatter(a, numPerProc, MPI_FLOAT, local_a, numPerProc,  
    MPI_FLOAT, 0, MPI_COMM_WORLD);  
2 MPI_Scatter(b, numPerProc, MPI_FLOAT, local_b, numPerProc,  
    MPI_FLOAT, 0, MPI_COMM_WORLD);
```

Each process then calculates its own local dot product:

```
1 float local_dot=0.0;  
2 for( i=0; i<numPerProc; i++ )  
3     local_dot += local_a[i]*local_b[i];
```

---

<sup>1</sup>This step is the same as for vector addition; cf. Lecture 9.

## MPI\_Reduce()

# Assignment Project Exam Help

The MPI standard supports reduction through `MPI_Reduce()`:

```
1 float dot;  
2 MPI_Reduce(&local_dot, &dot, 1, MPI_FLOAT, MPI_SUM, 0,  
    MPI_COMM_WORLD);
```

- Binary operator specified (`MPI_SUM`).
- Applied to `local_dot` on all processes.
- Reduced to `dot` on rank 0 (the 6<sup>th</sup> argument).
- Other operations are supported, e.g. `MPI_PROD`, `MPI_MAX`, `MPI_MIN`, logical and binary boolean operators.
- Implementation opaque to the user, but *should* be optimised for the system on which it is installed.

<https://powcoder.com>  
Add WeChat powcoder

## Efficient parallel reduction

# Assignment Project Exam Help

How OpenMP and MPI implement reduction is not specified by their respective standards.

- Allows **optimisation** for specific hardware architectures.

<https://powcoder.com>

Usually best to use the support as provided, but sometimes useful to consider possible implementation details to help understand performance and identify potential issues.

Add WeChat powcoder

Parallel reduction starts after each of **processing units** (threads, processes) have completed their **local reduction**.

- That is, calculated the **partial sums** of all the data each processing unit is 'responsible' for.

## Binary trees

# Assignment Project Exam Help

The most common method to implement parallel reduction is with a **binary tree**:

- One 'leaf' node for each processing unit.
- For  $p$  processing units, there are  $\log_2(p)$  levels<sup>1</sup>
- Perform calculations **in parallel** at each level.
- Reduction time is then  $\mathcal{O}(\log_2(p))$ , which is **much** faster than  $\mathcal{O}(p)$  for large  $p$ .

---


<sup>1</sup>If  $p$  is not a power of 2, round up.



## Binary tree: Example 1

Not all binary trees are valid for all binary operators  $\otimes$


For instance, this version requires that  $\otimes$  be **associative**:



The **indexing**, *i.e.* which processing units are performing the operations at each level, can be performed using bitwise arithmetic.

## Binary tree: Example 2

For this example,  $\otimes$  must also be commutative:



<https://powcoder.com>

Add WeChat powcoder

Indexing is easier than the previous example:

- In the first level, units 0 to  $p/2$  perform the operations.
- In the next level, units 0 to  $p/4$  perform the operations.
- ...


## Synchronisation between levels

Assignment Project Exam Help

Note we must ensure each level's calculations have been **completed** before continuing to the **next** level.

This example, where units 3, 6 and 7 are delayed, would result in at best an incorrect calculation, and at worst **deadlock**:

<https://powcoder.com>  
Add WeChat powcoder



## Barriers

# Assignment Project Exam Help

Most parallel APIs provide a means to synchronise all processing units at a specific point in code.

- Often called **barriers**.

<https://powcoder.com>

For instance, in OpenMP (in a parallel region):

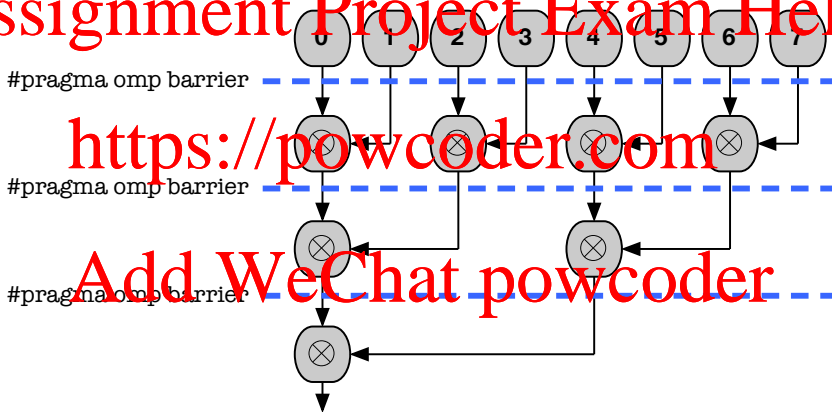
```
1 #pragma omp barrier
```

Add WeChat powcoder

- No processing unit (*i.e.* thread) will proceed past the barrier command until **all** units have reached it.

## Barrier synchronisation in a binary tree

Assignment Project Exam Help



<https://powcoder.com>

Add WeChat powcoder

## Synchronisation in MPI

# Assignment Project Exam Help

MPI also provides a barrier operation:

```
1 MPI_Barrier( MPI_COMM_WORLD );
```

However, there is usually no need as the necessary synchronisation can be achieved using **blocking communication**.

- `MPI_Send()`, `MPI_Recv()` will not return until message has been sent or received.
- Provides the necessary synchronisation between pairs of processes.

<https://powcoder.com>  
Add WeChat powcoder

## Binary trees in collective communication

# Assignment Project Exam Help

Note that `MPI_Reduce()` is a **collective communication**:

- Must be called by **all** ranks.

<https://powcoder.com>

The binary tree pattern is typically used for all collective communication.

- Communication time  $t_{\text{comm}} = \mathcal{O}(\log_2(p))$ .
- Faster than the  $\mathcal{O}(p)$  for a loop of send-and-receives.
- 'Inverted' in the case of `MPI_Bcast()` and `MPI_Scatter()`.

## Summary and next lecture

# Assignment Project Exam Help

Today we have looked at **parallel reduction**:

- Supported by most libraries, including OpenMPI and MPI.
- Typically implemented as a **binary tree**.
- Famous example was Google's **MapReduce**.
- In MPI, the necessary synchronisation provided by using **blocking communication**.

Add WeChat powcoder

Next time we will look at **non-blocking**, or **asynchronous**, communication.