

Assignment Project Exam Help

XJCO3221 Parallel Computation

<https://powcoder.com>

Peter Jimack

University of Leeds
Add WeChat powcoder

Lecture 13: Load balancing

Previous lectures

Assignment Project Exam Help

Several times in this module we have mentioned the concept of **load balancing**:

- Poor load balancing results in processing units spending time **idle**.
- Usually realised when **synchronising** threads or processes.
- First encountered for the Mandelbrot set generator [Lecture 3].
- Important for parallel performance for **all** architectures — shared and distributed memory CPU, and GPU.

Today's lecture

Assignment Project Exam Help

Today we will look at load balancing more closely, and how to reduce its performance penalty:

- Return to the example of the **Mandelbrot set** generator, this time in MPI.
- Understand how heterogeneity in the problem results in poor load balancing.
- See how a **task scheduler** can improve load balancing at **runtime**.
- Go through a concrete example of a **work pool**.

The Mandelbrot set (c.f. Lecture 3)

Code on Minerva: `Mandelbrot_MPI.c` plus makefile

Assignment Project Exam Help

- Domain is $-2 \leq x \leq 2$ and $-2 \leq y \leq 2$.
- Calculation performed **iteratively** for each (x, y) .
- Pixel coloured according to the number of iterations.
- Here, the **black region** corresponds to a **high number of iterations**.
- **No upper bound** - some points will iterate **indefinitely** if allowed.



Strip partitioning

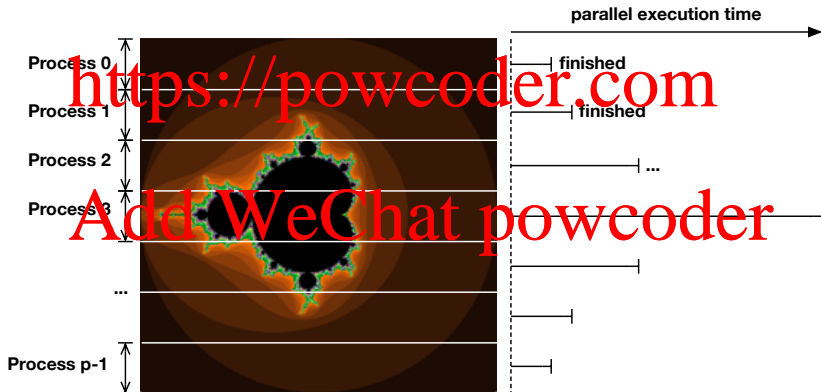
Partition the domain [cf. last lecture] into horizontal strips¹.



¹Equivalent results for partitioning into vertical strips, or blocks.

Load imbalance

Because some pixels take longer to calculate the colour than others, the load is **unevenly distributed** across the processes:



Load balancing

Assignment Project Exam Help

- Parallel execution time determined by the **last** processing unit to finish.
- Poor load balancing results in significant **idle time** for at least one process/thread.
- Inefficient use of available resources.

<https://powcoder.com>

Definition

The goal of **load balancing** is for each **processing unit** (thread or process) to perform a **similar volume of computations**, and therefore finish at roughly the same time.

Up until now most problems we have encountered have been naturally load balanced

Assignment Project Exam Help

For example, for vector addition between two n -vectors, assigning each processing unit to equal numbers of vector elements results in good load balancing.

- Each unit performs n/p additions.

Add WeChat powcoder

Note that the Mandelbrot set is a **map**, i.e. an **embarrassingly parallel problem** (since there are no data dependencies).

- *Still* a challenge to attain good performance.

Static load balancing

Assignment Project Exam Help

Definition

Sometimes it is possible to determine (approximately) equal loads **at compile time**. This is known as **static load balancing**.

<https://powcoder.com>

For the Mandelbrot set example, we could assign **larger** domains to regions where the calculations should be **fast**.

- Should improve load balancing.

Add WeChat powcoder

However, an **exact** expression is not available. Therefore any such **heuristic** can only achieve **approximate** load balancing.

Static load balancing (*ideal case*)

Assignment Project Exam Help



<https://powcoder.com>

Add WeChat powcoder

Dynamic load balancing

Assignment Project Exam Help

Definition

Dynamic load balancing is performed **at runtime**. No *a priori* knowledge of computation times is required.


<https://powcoder.com>

Basic idea:


- 1 Break the problem down into small **independent tasks**.
- 2 Each processing unit performs **one** task at a time.
- 3 When it is complete, it starts/is assigned another task.
- 4 Repeat 3 until all tasks are complete.

Add WeChat powcoder


Full problem broken
down into tasks:




Tasks performed sequentially on
different threads/processes:



Processing unit 1



Processing unit 2



Processing unit 3

scheduler

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Functional or task parallelism

Assignment Project Exam Help

Up to now, we have largely considered parallelising the same operation to a (large) data set.

- Known as **data parallelism**.
- When performed in a loop, can also be termed **loop parallelism**.

Now we are parallelising a number of **tasks**.

- Called **task parallelism** or **functional parallelism**.
- Be warned that these terms are sometimes used to refer to slightly different concepts.
- More on task parallelism in Lecture 19.

Work pools

Assignment Project Exam Help

- Often part of the parallel/concurrency runtime system.
- Introduces a (small) overhead.
- Various algorithms exist for efficient task scheduling.

To understand the role of a scheduler, we will look at a simple scheduler implemented in MPI - a **centralised work pool**.

- One process (usually rank 0) performs the scheduling - this is the **main** process¹.
- Remaining processes action the tasks - the **workers**¹.

¹You may see 'master' (for main) and 'slaves' (for workers) in the literature.

Worker pseudocode

Function workerProcess() in Mandelbrot_MPI.c

```
1 initialise(); // Including MPI_Init().
2
3 while( true )
4 {
5     // Wait for message from the main (rank 0).
6     MPI_Recv( message, ... );
7
8     // Is this a termination request?
9     if( message==TERMINATE ) break;
10
11    // Else perform calculation and send back to rank 0.
12    result = actionTask( message );
13    MPI_Send( result, ... );
14 }
15
16 finalise(); // Including MPI_Finalize().
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Main process pseudocode (1)

Function mainProcess() in Mandelbrot_MPI.c

Assignment Project Exam Help

```
1 initialiseAndOpenWindow();  
2  
3 // Initialise variable that tracks progress.  
4 int numActive=0;  
5  
6 // Send initial request to each worker.  
7 for( p=1; p<numProcs; p++ )  
8 {  
9     MPI_Send(task, ..., p, ...);  
10    numActive++;  
11 }
```

<https://powcoder.com>

Add WeChat powcoder

For this Mandelbrot example, each task is a **row of pixel colours to be calculated**.

- Keep track with an incrementing variable row.

Main process pseudocode (2)

Function idle() in Mandelbrot_MPI.c

```
1 while( numActive > 0 )
2 {
3     // Get result from ANY worker process.
4     MPI_Recv(result, ..., MPI_ANY_SOURCE, ..., &status);
5     numActive--;
6
7     // Send request IMMEDIATELY to the SAME worker.
8     if( !finished )
9     {
10        MPI_Send(task, ..., status.MPI_SOURCE, ...);
11        numActive++;
12    }
13
14    // Action the message.
15    actionResult( result );
16 }
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Main process pseudocode (3)

Function `idle()` in `Mandelbrot_MPI.c`

Assignment Project Exam Help

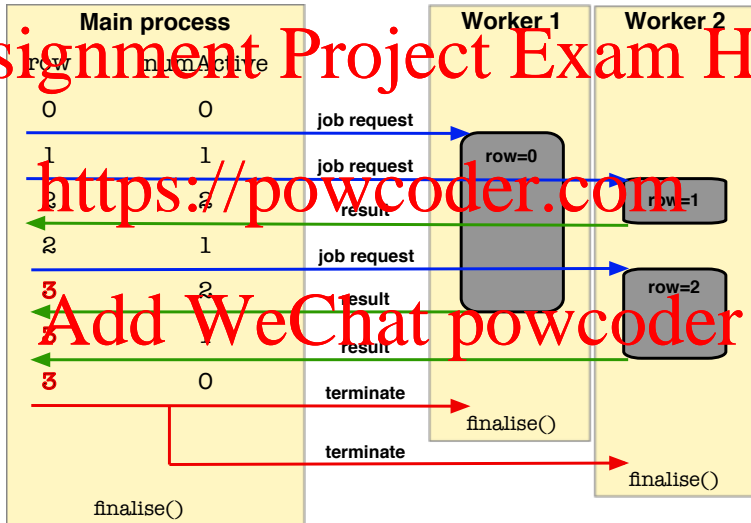
```
1 // Tell all workers to terminate.  
2 for( p=1; p<numProcs; p++ )  
3     MPI_Send(TERMINATE,...);  
4  
5 finaliseAndCloseWindow();
```

<https://powcoder.com>

- `MPI_ANY_SOURCE` in place of source in `MPI_Recv()` receives a message from any process.
- Used `status.MPI_SOURCE` to recover the rank of the sending process.
- Send next request **before** the (potentially slow) call to `actionResult()`.

Add WeChat powcoder

Example: 3 rows and 2 workers



Modern schedulers

Assignment Project Exam Help

There are many more types of **work pool**, such as those with no 'main' process (*decentralised work pools*)¹.

A common approach is to use **dequeues** or *double ended queues*:

- Each processing unit maintains its own **deque** of tasks.
- Performs tasks sequentially, starting from the front.
- Once the deque is empty, 'steals' a task from the **back** of a randomly selected 'victim' (**work stealing**).

¹Wilkinson and Allen, *Parallel Programming* (Pearson, 2005).

OpenMP scheduler

Assignment Project Exam Help

OpenMP can also schedule loops using its `schedule` clause:

```
1 #pragma omp parallel for schedule(dynamic, chunk)
2 for( i=0; i<n; i++ ) { ... }
```

This breaks down the loop into 'chunks' of size `chunk` at runtime.

Can also be used for **static** scheduling:

```
1 #pragma omp parallel for schedule(static, chunk)
```

There is also a **guided** option that decreases the chunk size exponentially **at runtime** to the final value `chunk`:

```
1 #pragma omp parallel for schedule(guided, chunk)
```

In all cases, `chunk` is optional and defaults to 1.

MIMD at last!

Assignment Project Exam Help

Up until today we have mostly performed the **same** calculations on each processing unit.

- SIMD (Single Instruction Multiple Data) **software** ...
- .. on MIMD (Multiple Instruction Multiple Data) **hardware**.

Today is the first clear¹ example where we have implemented the MIMD pattern **in software**.

- The **main** process perform entirely different calculations to **workers** (*division of labour*).

¹Ignoring trivial cases like e.g. rank 0 distributing global arrays.

Summary of distributed memory systems

Lec.	Content	Key points
8	Architectures and MPI	Clusters and supercomputers; interconnect network; starting with MPI.
9	Point-to-point communication	Blocking send and receives; buffering; deadlock for cyclic communication.
10	Data reorganization	Scatter and gather; t_{comm} ; collective communication in MPI.
11	Reduction	Binary trees; OpenMP and MPI.
12	Asynchronous communication	Non-blocking send and receives; domain partitioning and ghost cells.
13	Load balancing	Task parallelism; schedulers; work pools.

Next lecture we start looking at programming **general purpose graphics processing units** or GPGPUs.