

Assignment Project Exam Help

XJCO3221 Parallel Computation

<https://powcoder.com>

Peter Jimack

University of Leeds
Add WeChat powcoder

Lecture 18: Atomic operations

Previous lectures

Assignment Project Exam Help

Whenever multiple processing units had read-write access to the same memory location, there are potential **data races**:

- If at least one unit **writes** to the memory [Lecture 5].
- Can solve using **critical regions** guarded by **locks/mutexes** [Lectures 6 and 7].
- Single instructions can also be made performed **atomically** [Lecture 6].

Add WeChat powcoder

For instance, in OpenMP an atomic instruction looks like:

```
1 #pragma omp atomic
2 count++;
```

Today's lecture

Assignment Project Exam Help

GPUs also have memory accessible by multiple work items / threads:

- **Global** memory, accessible to all work groups.
- **Local** memory, only accessible within a work group

There is therefore potential data races and the need for **synchronisation**.

Add WeChat powcoder

Today we will see how GPUs support **atomic operations** in a similar way to a shared memory CPU.

- Also consider an atomic **compare and swap**.

Atomic operations

Assignment Project Exam Help

Definition

An **atomic operation** is one that is completed **without interruption** by any other processing unit.

<https://powcoder.com>

- Usually restricted to simple **arithmetic operations** (addition, subtraction etc.)
- Implemented by a combination of compiler and hardware.
- Typically a (much) **smaller performance penalty** than using locks/mutexes etc.

Add WeChat powcoder

Load, compute, and store

Consider the following line:

```
1 x -= 2;
```

Even this single instruction performs **three** sub-operations:

- 1 **Loads** the value of `x`.
- 2 Performs the **computation** (i.e. subtracts 2).
- 3 **Stores** the updated value.

Two or more processing units might interfere with each other, resulting in a different result to the **serial** equivalent.

This could not happen if the operation was **atomic**.

Example

Assignment Project Exam Help

Suppose $x=10$ initially, and two processing units A and B both subtract 2 from x . Depending on the scheduler, this may happen:

- 1 A loads the value of x as 10.
- 2 B loads the value of x as 10.
- 3 A performs its computation: $10 - 2 = 8$.
- 4 A stores 8 to memory.
- 5 B performs its computation: $10 - 2 = 8$.
- 6 B stores 8 to memory.

The result is $x = 8$, rather than $x = 6$ as expected.

Constructing a histogram on a GPU

Code on Minerva: `histogram.c`, `histogram.cl`, `helper.h`

Assignment Project Exam Help

Have an array of integers in the range 0 to `maxValue-1` inclusive; want the **histogram** showing the frequency of each value.

- ① Memory allocated on the **host** and on the **device**, for both the data and the histogram.
 - `data`, `hist` on the host.
 - `device_data`, `device_hist` on the device.
- ② Both initialised on the host and copied to the device.
- ③ Build, compile and enqueue a kernel to construct the device histogram.
 - **One work item per data element**, e.g. `data[i]`.
- ④ Copy the histogram back to the host using `clEnqueueReadBuffer()`.

<https://powcoder.com>

Add WeChat powcoder

Kernel 1: Direct to global; no atomics

Assignment Project Exam Help

```
1 // kernel
2 void histogramNoAtomic(
3     __global int *device_hist,
4     __global int *device_data,
5     int max_value )
6 {
7     int gid = get_global_id(0);
8
9     // Data value.
10    int val = device_data[gid];
11
12    // Check range before updating.
13    if( val >= 0 && val < max_value )
14        device_hist[val]++;
15 }
```

<https://powcoder.com>

Add WeChat powcoder

Kernel 2: Direct to global; atomic.

Code fails because the update of `device_hist` is not atomic

```
1 device_hist[val]++;           // Load, compute, store.
```

Many additions to the histogram are lost, resulting in lower totals.

In OpenCL, can make this atomic by using `atomic_inc()`:

```
1 __kernel
2 void histogramAtomic( ... )
3 {
4     ... // as before
5     if( val >= 0 && val < maxValue )
6         atomic_inc( &(amp;device_hist[val]) );
7 }
```

This now works as expected.

Atomic operations in OpenCL

Assignment Project Exam Help

There are many atomic operations in OpenCL¹:

| | |
|---|--|
| <code>atomic_inc</code> , <code>atomic_dec</code> | Increment, decrement. |
| <code>atomic_add</code> , <code>atomic_sub</code> | Addition, subtraction. |
| <code>atomic_min</code> , <code>atomic_max</code> | Smaller or larger of two arguments. |
| <code>atomic_and</code> , <code>atomic_or</code> , <code>atomic_xor</code> | Bitwise operations. |
| <code>atomic_xchg</code> , <code>atomic_cmpxchg</code> | Exchange, compare and exchange (<i>see later</i>). |

¹Similar in CUDA, *i.e.* `atomicAdd()`, `atomicInc()` etc.

Optimising with local memory

Assignment Project Exam Help

Having a **single** histogram in global memory is not efficient:

- Potential for **very many** work items attempting to access the **same** global memory location **almost simultaneously**.

<https://powcoder.com>

More efficient for each **work group** to calculate its own histogram **in local memory**, then update the global histogram **at the end**.

- Fewer **competing** work items for the local histogram.
- Local memory is faster anyway.

Add WeChat powcoder

Aside: Could use a similar strategy for a *multi-threaded CPU* (i.e. each *thread* constructs its own histogram).

Kernel 3: Local histogram (1)

```
1 _kernel
2 void localHistogram(..., __local int *local_hist)
3 {
4     int
5     i,
6     lid = get_local_id(0),
7     gid = get_global_id(0),
8     size = get_local_size(0);
9
10    // Clear the histogram
11    for(i=lid; i<maxValue; i+=size)
12        local_hist[i] = 0;
13
14    // Ensure histogram fully initialised.
15    barrier(CLK_LOCAL_MEM_FENCE);
16
17    // (cont'd next slide).
```

Kernel 3: Local histogram (2)

```
1 // (from previous slide)
2
3 // Add to the local histogram.
4 int val = device_data[gid];
5 if (val >= 0 && val < maxValue )
6     atomic_inc( &(amp;local_hist[val]) );
7
8 // Ensure local histogram calculation complete
9 // before moving on.
10 barrier(CLK_LOCAL_MEM_FENCE);
11
12 // Atomic add the local histogram to the global one.
13 for( i=lid; i<maxValue; i+=size )
14     atomic_add( &(amp;device_hist[i]), local_hist[i] );
15 }
```

You should see a performance improvement using this method.

Could have had **one** work item in each group initialise and update the **entire** local histogram, e.g.:

```
1 if( gid==0 )  
2   for( i=0; i<maxValue; i++ )  
3     atomic_add( &(amp;device_hist[i]), local_hist[i] );
```

This would work, but would be slower (check).

Instead use as many work items as possible:

```
1 for( i=gid; i<maxValue; i+=size )  
2   atomic_add( &device_hist[i], local_hist[i] );
```

- Each i in the range realised by **exactly** one work item.
- Spans full range even if $\text{size} < \text{maxValue}$.

Atomic exchange and compare-and-exchange

Assignment Project Exam Help

```
int atomic_exchg(int *p, int val).
```

- 1 Sets old=*p.
- 2 Sets *p=val.
- 3 Returns old.

<https://powcoder.com>

```
int atomic_cmpxchg(int *p, int cmp, int val):
```

- 1 Sets old=*p.
- 2 Sets *p=val if *p==cmp, otherwise does not change.
- 3 Returns old.

In both cases, p can be in local or global memory, and the data type can be int, unsigned int or float.

Uses of compare and exchange

Many low-level parallel frameworks provide similar functionality, sometimes referred to as CAS = Compare And Swap:

- `atomicExch()`, `atomicCAS()` in CUDA;
- `std::atomic::compare_exchange_weak()`,
`std::atomic::exchange()` in C++11.

Common uses include:

- 1 Implementing a **lock** or **mutex**.
- 2 **Lock-free** implementations.

Examples below are for OpenCL, but just as relevant for CUDA and multi-core CPUs.

Spinlock

Assignment Project Exam Help

- Suppose an int variable `lock` was accessible to multiple threads.
- `lock` takes **two** values, 0 and 1.
- Take 0 to be **unlocked**, 1 to be **locked**.

A simple 'spinlock' can be implemented as follows:

```
1 int lock;           // 0 or 1. Accessible by all threads.  
2 ...  
3 while (atomic_compare_exchange(&lock, 0, 1) != 1);
```

- Infinite while loop, until `lock==0`.
- Then sets `lock=1` and continues past line 3.
- Does all this **atomically**.

Why atomic?

Assignment Project Exam Help

Consider what could happen **without** atomicity:

```
1 while( lock==1 );  
2 lock = 1;  
3 ...
```

<https://powcoder.com>

- ① One thread / work item sees `lock==0` and continues to line 2.
- ② A second thread **also** sees `lock==0` and proceeds to line 2, before the first thread sets `lock=1`.
- ③ The first thread now sets `lock=1`.
- ④ The second thread **also** sets `lock=1`.
- ⑤ **Both** continue to line 3!

Spinlocks vs. locks/mutexes:

- **Pro:** Spinlocks **faster**, as do not put the thread to sleep.
- **Con:** Spinlocks waste CPU/GPU cycles.

Spinlocks on GPUs:

Note Lock could be in **global** memory

- Accessible to work items from **different** work groups.
- May seem this can be used to **synchronise between** groups.

But recall the warning from Lecture 17

Cannot guarantee all work groups are active on the device **at the same time** (as some may be queued), so this not a **robust** synchronisation mechanism.

Lock-free data structures

Atomic compare and exchange can also be used to implement **thread-safe** access to data structures **without requiring locks**, and the associated overhead.

<https://powcoder.com>
Such **lock free** data structures, if they can be achieved, are desirable for good parallel performance.

Example: **Appending** an item to a singly linked list.

- Need to ensure old and new head nodes updated **together**
- Use `atomic_cmpxchg()` in an infinite loop¹.

¹McCool *et al.*, *Structured parallel programming* (Morgan-Kaufman, 2012).

Basic idea of a lock-free data linked list

```
1 struct node { ... }; // Some data structure
2 node *head; // Head of list.
3 void prependToList( node *a ) // 'a' becomes head.
4 {
5     while(true)
6     {
7         // Take a copy of current head pointer.
8         node *b = head;
9
10        // Link to the node being added.
11        a->next = b;
12
13        // Only update head if not just changed by another
14        // work item/thread; else try again from line 6.
15        if( atomic_cmpxchg(head,b,a)==b ) break;
16    }
17 }
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```
if( atomic_cmpxchg(head,b,a)==b ) break;
```

Assignment Project Exam Help

If only a single thread was involved:

- 1 old=*head, i.e. old==b, the first item in the list.
- 2 **Compare-exchange:** *head==b, so changes *head to a.
- 3 atomic_cmpxchg returns b, so will break from while loop.

This is the expected behaviour.

However, in a multi-threaded context:

- 1 Another thread may change *head from b before line 15.
- 2 Since *head!=b, will **not** change it.
- 3 Will return some value !=b, so will try again.

Summary and next lecture

Assignment Project Exam Help

This lecture we have revisited **atomic operations** with an emphasis on GPUs:

- Atomics used to ensure correct updates of memory accessible by multiple work items.
- Atomic **compare and exchange** can be used to implement a spinlock, lock-free data structures, *etc.*

Add WeChat powcoder

The next lecture is the last on GPU programming when we will look at events and **task parallelism**.